



Figure 3. Search behaviour



(3B) MCTS root-action visits

MCTS.MCTSSolver (ff95b-mcts)
state 65ecc28ddcd8e0c5
config b29cbf96cc8fb389 seed 2026
nodes 289 depth 2
root 36 child nodes, 7 distinct actions, 36 visits
per action mean 5.14 median 2 max 24 once-only 42.9%
selected SLOW_STRAIGHT
value mean of backed-up returns (MCTS.jl updates $q \leftarrow (return - q)/n$, i.e. an incremental sample mean; the first value is seeded by $init_Q$, not by a rollout)



(3D) DPW sampled action plane




Figure 3. Search behaviour at a single frozen state. Both planners searched from the same state (fingerprint "65ecc28ddcd8e0c5"), at a matched budget. (3A, 3B) MCTS expands 7 distinct root actions with 5.14 visits per action on average (median 2, max 24). (3C, 3D) DPW expands 24 distinct continuous root actions with mean 1.46 (median 1, max 3); only sampled actions are drawn, since a smoothed surface would imply combinations the planner never tried. Both planners leave many candidates lightly explored, but MCTS concentrates substantially more evidence on its most-preferred discrete action, whereas DPW spreads the same budget across many continuous actions.

figure_id = Figure 3
panel_ids = 3A,3B,3C,3D
layout_version = ff9b.1
3A.solver = MCTS.MCTSSolver
3A.state = 65ecc28ddcd8e0c5
3A.root_actions = 7
3A.mean_visits = 5.14
3A.median_visits = 2
3A.max_visits = 24
3B.total_root_visits = 36
3C.solver = MCTS.DPWSolver
3C.state = 65ecc28ddcd8e0c5
3C.root_actions = 24
3C.mean_visits = 1.46
3C.median_visits = 1
3C.max_visits = 3
3D.sampled_actions = 24
state_fingerprint = 65ecc28ddcd8e0c5
mcts.root_actions = 7
dpw.root_actions = 24